

Clab-2 Report

ENGN6528

Shreya Chawla

u7195872

25/04/2021

Task-1 Tasks Harris Corner Detector. (6 marks):

For Python users:

1. The completed code for Task 1 is in “harris.py” in “Task 1” folder. All the results and visualizations obtained are saved in “Task 1” folder as well.
2. The “harris.py” python script consists of my_harris() function, non_maximum_suppression() function, plot_corner_img() function to plot corners, and functions originally defined in Lab2 code provided - conv2() and fspecial() functions.

The following is the code for “Task: Complete the Harris Cornerness” (Fig 1). The arguments are grayscale image “im”, “sigma” used to find window ‘g’, threshold “thresh”, “k” parameter in cornerness formula. It returns an array containing corners’ x, y coordinates and response value.

```
42
43 '''
44 my_harris : Function to detect Harris corners on the given image
45 '''
46 def my_harris(im, sigma, thresh, k):
47     # Find x and y gradient
48     dx = np.array([[ -1,  0,  1], [ -1,  0,  1], [ -1,  0,  1]])
49     dy = dx.transpose() # 'dy' is transpose of 'dx'
50     Ix = conv2(im, dx)
51     Iy = conv2(im, dy)
52
53     # Find corners in image
54     # 'g' is the window function which will be used to find R (response). It is a Gaussian window function.
55     # Dimensions of 'g' is neighborhood to be considered.
56     g = fspecial('gaussian', (max(1, np.floor(3 * sigma) * 2 + 1), max(1, np.floor(3 * sigma) * 2 + 1)), sigma)
57     Iy2 = conv2(np.power(Iy, 2), g) # y and y
58     Ix2 = conv2(np.power(Ix, 2), g) # x and x
59     Ixy = conv2(Ix * Iy, g) # x and y
60
61     # Find if interest point using cornerness
62     # Determinant
63     det = Ix2 * Iy2 - Ixy * Ixy
64     # Trace
65     trace = Ix2 + Iy2
66     # R is the cornerness or response
67     R = det - k * trace * trace
68
69     # Threshold based on max response value
70     thresh *= np.amax(R) # Threshold for an optimal value, it may vary depending on the image.
71
72     # List of corners detected
73     corners = []
74
75     # If response > threshold then it is a corner else skip
76     for row, response in enumerate(R):
77         for col, r in enumerate(response):
78             if r > thresh:
79                 corners.append([row, col, r])
80     print('\nNumber of corners detected by my Harris corner detector = ', len(corners))
81     return corners
82
```

Fig 1: Task - Complete the Harris Cornerness

The following is the code for “Task: Perform non-maximum thresholding and return Nx2 matrix of x and y coordinates”. (Fig 2)

The arguments for this function are array of “corners” (a 3-D array containing x, y coordinates of corners and their cornerness score) and distance “dist”, which is the neighborhood on which suppression is performed.

```

83
84 '''
85 non_maximum_suppression : Function performs non-maximum suppression corners at distance dist
86 '''
87 def non_maximum_suppression(corners, dist):
88     # If no corners found then return that empty list
89     if len(corners) == 0:
90         return corners
91
92     # Sort corners based on their cornerness value
93     corners = sorted(corners, key=lambda c: c[2], reverse=True)
94     # List of corners that are not suppressed
95     chosen_corners = list()
96     chosen_corners.append(corners[0][: -1])
97
98     # Compare corners to see if there are more corners in vicinity of the current corner
99     # If not in the neighborhood of dist distance then that corner is chosen
100    for corner in corners:
101        for chosen in chosen_corners:
102            if abs(corner[0] - chosen[0]) < dist and abs(corner[1] - chosen[1]) < dist:
103                break
104        else:
105            chosen_corners.append(corner[: -1])
106    print('Number of corners detected after non maximum suppression = ', len(chosen_corners))
107
108    return chosen_corners
109

```

Fig 2: Task - Perform non-maximum thresholding and return Nx2 matrix of x and y coordinates

3. The following comments in code corresponding to block #5 have been added as can be seen in Fig 3. (It is a part of Fig 1 according to my function definition).

```

52
53 # Find corners in image
54 # 'g' is the window function which will be used to find R (response). It is a Gaussian window function.
55 # Dimensions of 'g' is neighborhood to be considered.
56 g = fspecial((max(1, np.floor(3 * sigma) * 2 + 1), max(1, np.floor(3 * sigma) * 2 + 1)), sigma)
57

```

Fig 3: Comments on the codes in block #5

Comments added to code corresponding to block #7 are in Fig 1 and Fig 7 and explained in 2nd part of Task 1.

The following comments have been added to the fspecial function which has been used in block #5. (Fig 4)

```


30 def fspecial(shape=(3, 3), sigma=0.5): # Default values of shape of window and sigma for Gaussian
31     m, n = [(ss - 1.) / 2. for ss in shape]
32     y, x = np.ogrid[-m:m + 1, -n:n + 1] # Creates a multidimensional meshgrid
33     h = np.exp(-(x * x + y * y) / (2. * sigma * sigma)) # Fourier tranform of x
34     # finfo() gives machine limits for floating point types, eps is the difference between 1.0 and the next smallest
35     # representable float larger than 1.0. Thus np.finfo(h.dtype).eps would be the smallest float in magnitude.
36     h[h < np.finfo(h.dtype).eps * h.max()] = 0
37     sumh = h.sum() # Sum of h array (which are > 0 in previous step)
38     if sumh != 0:
39         h /= sumh # Normalize h if sum > 0
40     return h

```



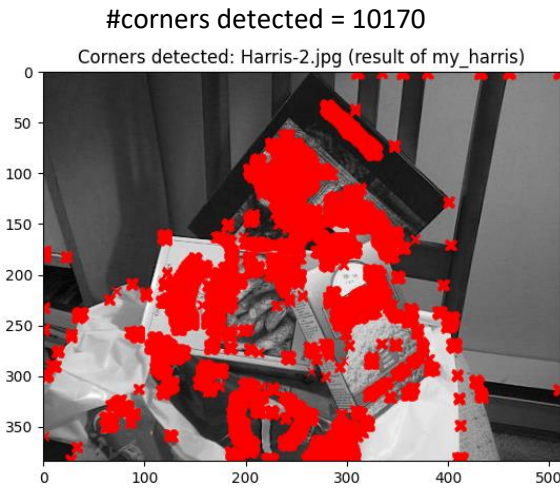

Fig 4: Comments on the codes in block #3 used in block #5

4. The results observed for the 4 test images (Harris-[1,2,3,4].jpg) are shown in Table 1 below. Red crosses have been used to display corners. Also, zoomed images for better observing corners detected in “Harris-3.jpg” and “Harris-4.jpg” are Fig 5 and Fig 6 respectively.


Result of my_harris() (Before suppression)



Result of non_maximum_suppression() (After suppression)




Corners detected: Harris-4.jpg (result of my_harris)



#corners detected = 4030

Corners detected: Harris-4.jpg (after suppression)



#corners detected = 78

Table 1: The left image corresponds to corners obtained after applying my_harris function (without suppression) and the right image corresponds to corners obtained after non_maximum_suppression function is performed.

The my_harris and non_maximun_suppression functions are able to detect many corners in the image as clearly visible in the zoomed images below. (Fig 5 and Fig 6).

Corners detected: Harris-3.jpg (after suppression)




Fig 5: Zoomed Harris corners detected in Harris-3.jpg

Corners detected: Harris-4.jpg (after suppression)




Fig 6: Zoomed Harris corners detected in Harris-4.jpg

5. The result of in-built `cv2.cornerHarris()` function on “Harris-1.jpg” is shown in Fig 7. Upon comparison, although the results are similar, the inbuilt function detects more corners than `my_harris()`. This could be due to different values of parameters used in both the functions although the arguments for both functions have same value.

Harris corner is affected by scale of image and is also dependent on the rotation of image. The threshold and the value of ‘k’ in cornerness response score are also some factors that affect the performance of Harris corner detection.

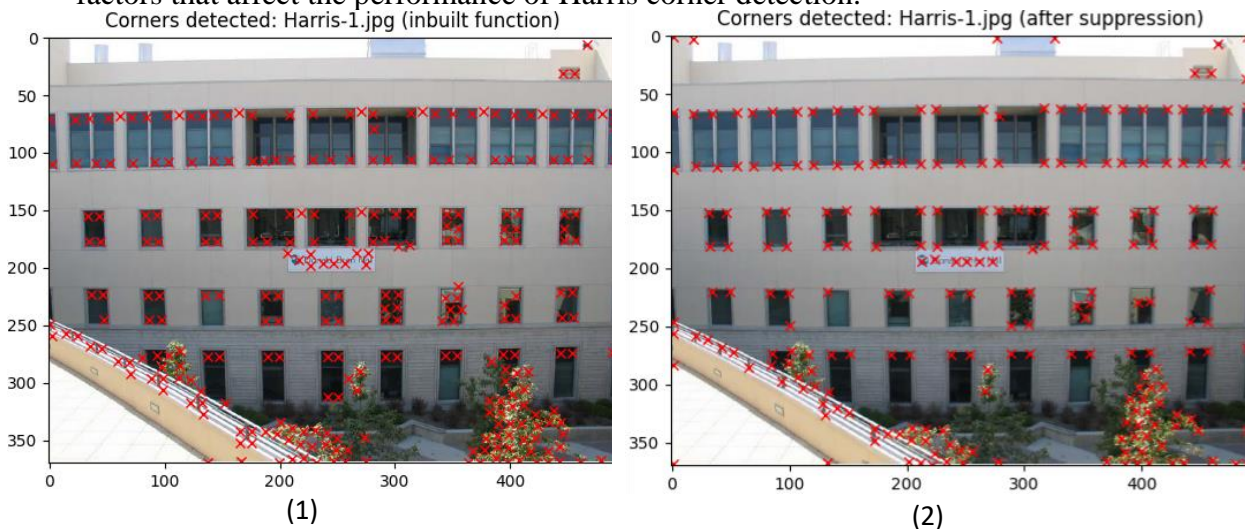


Fig 7: Comparison between Harris corners detected in “Harris-1.jpg” using `cv2`’s inbuilt function (1) (left) and detected by `my_harris` (2) (right)

6. In 'Harris-5.jpg', we cannot get any corner as the corner response score (R) is less than threshold for all windows in image. That is because there is an edge in this image (i.e. $R < 0$) which cannot be detected by Harris corner technique. Fig 8 shows the results visualized below.

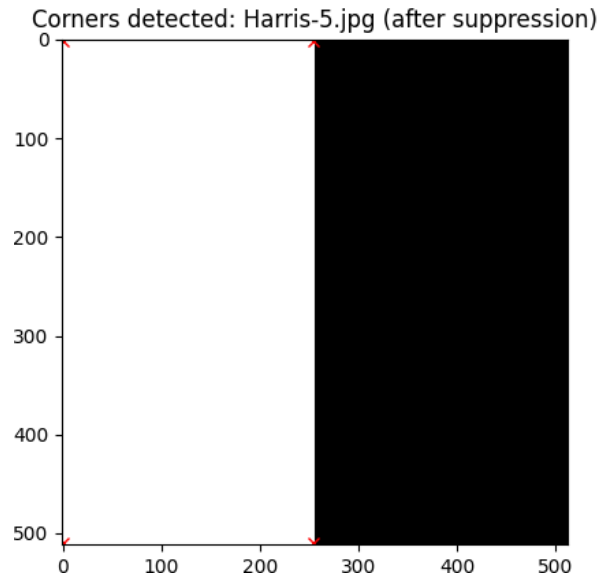


Fig 8: Harris corners for "Harris-5.jpg"

7. The visualization of results for "Harris-6.jpg" are in Fig 9. "Harris-6.jpg" contains a lot of noise hence it gets a lot of false positives in detecting corner. A smoothening filter like Bilateral filter can smoothen the noise while keeping the corner (and edges) sharp. This can help improve the corner detection.

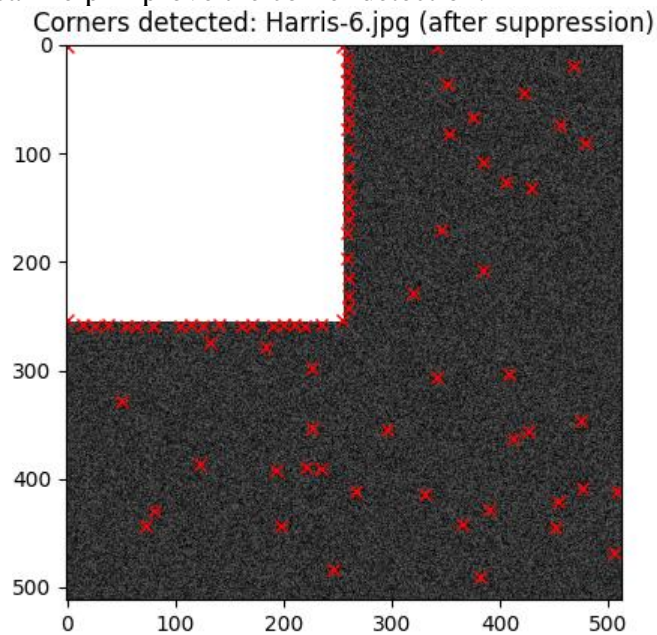


Fig 9: Harris corners for "Harris-6.jpg"

Task-2: K-Means Clustering and Color Image Segmentation. (5 marks)

In this task, you are asked to implement your own K-means clustering algorithm for colour image segmentation, and test it on the following **two images** as shown in Fig.1 (you can download them from wattle). Please first convert one image (in PNG-type (Portable Network Graphics)) of 16bit to 8bit image before the following process.



Fig. 1

1. The implementation of `my_kmeans()` is in Fig 10 below. The input is the data points to be processed (data) and the number of clusters (K), and the output is several clusters data (clusters).

```
6 '''
7 my_kmeans function: Performs K-means algorithm on data and returns the K clusters
8 '''
9 def my_kmeans(data, K):
10     start_time = time.time()
11     # Step 1 Initialize the K centroids randomly
12     centroids = np.random.random((K, data.shape[1]))
13     print('Initialized centroids to: ', centroids)
14     clusters = np.zeros_like(data) # Keeps track of which datapoint belongs to which cluster
15     diff = True # Boolean which keeps track of when to stop - reached local maximum
16
17     while diff:
18         new_centroids = np.zeros_like(centroids) # New centroids calculated are stored in this array
19
20         # Step 2 Find Euclidean distance (L2 norm) of each point from each centroid and store it in a dictionary
21         for i in range(data.shape[0]):
22             dist = dict() # Dictionary to store distance of each centroid from each point
23             for j in range(K):
24                 dist.update({j: np.linalg.norm(data[i] - centroids[j])})
25             # Step 3 Assign points to closest cluster (or centroid)
26             index = [key for key in dist if dist[key] == min(dist.values())]
27             if len(index) > 1: # If more than 1 centroids are at minimum distance
28                 index = index[0]
29             clusters[i] = centroids[index]
30
31         # Step 4 Compute the new centroid
32         for i in range(K):
33             cluster = np.where((clusters == centroids[i]).all(axis=1))[0]
34             if cluster.shape[0] > 0:
35                 new_centroids[i] = sum(data[cluster]) / cluster.shape[0]
36             else:
37                 new_centroids[i] = centroids[i]
38
39         # Step 5 If no difference between new and old set of centroids, then stop, else repeat step 2 to 4
40         if np.isclose(centroids, new_centroids, 1e-15).all():
41             diff = False
42             # print('Game over')
43         else:
44             print('diff = ', centroids - new_centroids)
45             centroids = new_centroids
46         end_time = time.time()
47         print('Time taken for the k means algorithm to run = ', end_time - start_time)
48     return clusters
```

Fig 10: Code for `my_kmeans` function

2. The images in LAB color space are displayed in Fig .




Fig 11: 'mandm.png' and 'pepper.png' in LAB color space

The boolean “with_coordinates” keeps track of whether to use 5D data or 3D. The 5-D vector that encodes: (1) L* - lightness of the color; (2) a* - the color position between red and green; (3) b* - the position between yellow and blue; (4) x, y - pixel coordinates. The code for the same is below in Fig 12.

```

154 with_coordinates = False # Boolean for how the segmentation is to be carried out (with or without pixel coordinates)
155 clusters_k = np.array([]) # Stores all the clusters returned from my_kmeans method (K-means)
156 clusters_k_pp = np.array([]) # Stores all the clusters returned from my_kmeans_pp method (K-means++)
157 if with_coordinates: # With pixel coordinates
158     # Create 5-D image data: L, A, B, X, Y
159     fiveD_data = list()
160     for i in range(im.shape[0]):
161         for j in range(im.shape[1]):
162             fiveD_data.append(np.append(np.append(im[i, j], np.array([i, j]))))
163     fiveD_data = np.array(fiveD_data)
164     fiveD_data = normalize(fiveD_data)
165     clusters_k = my_kmeans(fiveD_data, K)
166     # Take only first 3 values (L,A,B) to show image, drop x,y coordinate of centroid
167     clusters_k = np.reshape(clusters_k[:, :3], im.shape)
168     # Similarly for K-means++ clusters
169     clusters_k_pp = my_kmeans_pp(fiveD_data, K)
170     # Take only first 3 values (L,A,B) to show image, drop x,y coordinate of centroid
171     clusters_k_pp = np.reshape(clusters_k_pp[:, :3], im.shape)
172
173 else: # Without pixel coordinates
174     # Create 3-D image data: L, A, B
175     threeD_data = np.reshape(im, (im.shape[0] * im.shape[1], 3))
176     threeD_data = normalize(threeD_data)
177     clusters_k = my_kmeans(threeD_data, K)
178     clusters_k_pp = my_kmeans_pp(threeD_data, K)
179     clusters_k = clusters_k.reshape(im.shape)
180     clusters_k_pp = clusters_k_pp.reshape(im.shape)
181

```

Fig 11: Code for dealing with coordinates or without

3. The key steps for K-means++ algorithm are initialization and convergence. The convergence of K-means++ algorithm is performed the same way as K-means algorithm. They only differ in initialization. Let $D(x)$ denote the shortest distance from a data point in X to the closest center that was already chosen. The steps for the complete K-means++ algorithm are:

1. Take one centroids c_1 , uniformly sampled from X .
2. Take a new centroids c_i , choosing $x \in X$ with probability $D(x)^2 / \sum_{x \in X} D(x)^2$.

3. Repeat step 2 until the number of centroids is k.
4. For each $i \in \{1, \dots, k\}$, set the cluster C_i to be the set of points in X that are closer to c_i than they are to c_j for all $j \neq i$.
5. For each $i \in \{1, \dots, k\}$, c_i is set to be the center of mass of all points in cluster

$$C_i: c_i = 1/|C_i| \sum_{x \in C_i} x.$$

6. Repeat Steps 4 and 5 until the clusters C no longer change.

Steps 1 to 3 are the initialization steps which differ from K-means. The rest are the same steps.

The corresponding code is below (Fig 12)

```

49 """
50 my_kmeans_pp function: Performs K-means++ algorithm on data and returns the K clusters
51 """
52 def my_kmeans_pp(data, K):
53     # Step 1 Initialize the K centroids randomly
54     centroids = []
55     # 1.1) Choose the 1st centroid uniformly sampling from the data points in data
56     centroids.append(data[np.random.randint(data.shape[0])])
57     # 1.2) For the next K-1 centroids, we need to choose them as far from the 1st as possible
58     for c_id in range(K - 1):
59         dist = [] # Shortest distance between point and closest centroid (D(x))
60         for i in range(data.shape[0]):
61             point = data[i, :]
62             # For every point in X, choose the min distance between
63             # np.inf and L2 norm (or euclidean distance between points and the centroids)
64             d = np.inf
65             for j in range(len(centroids)):
66                 temp_dist = np.linalg.norm(point - centroids[j])
67                 d = min(d, temp_dist)
68             dist.append(d)
69
70     # Select data point with maximum distance from the 1st centroid as the next centroid
71     dist = np.array(dist)
72     next_centroid = data[np.argmax(dist)]
73     centroids.append(next_centroid)
74     # 1.3) Repeat 1.1 and 1.2 until there are K clusters
75
76 centroids = np.array(centroids)
77
78 print('Initialized centroids to: \n', centroids)
79 start_time = time.time() # Keep track of convergence time
80 clusters = np.zeros_like(data) # Keeps track of which datapoint belongs to which cluster
81 diff = True # Boolean which keeps track of when to stop - reached local maximum
82
83 while diff:
84     new_centroids = np.zeros_like(centroids) # New centroids calculated are stored in this array
85     # Step 2 Find Euclidean distance (L2 norm) of each point from each centroid and store it in a dictionary
86     for i in range(data.shape[0]):
87         dist = dict() # Dictionary to store distance of each centroid from each point
88         for j in range(K):
89             dist.update({j: np.linalg.norm(data[i] - centroids[j])})
90         # Step 3 Assign points to closest cluster (or centroid)
91         index = [key for key in dist if dist[key] == min(dist.values())]
92         if len(index) > 1: # If more than 1 centroids are at minimum distance
93             index = index[1]
94         clusters[i] = centroids[index]
95
96     # Step 4 Compute the new centroid
97     for i in range(K):
98         cluster = np.where((clusters == centroids[i]).all(axis=1))[0]
99         if cluster.shape[0] > 0:
100             new_centroids[i] = sum(np.abs(data[cluster])) / cluster.shape[0]
101         else:
102             new_centroids[i] = centroids[i]
103
104     # Step 5 If no difference between new and old set of centroids, then stop, else repeat step 2 to 4
105     if (centroids - new_centroids).sum() == 0:
106         diff = False
107     else:
108         centroids = new_centroids
109 end_time = time.time()
110 print('Time taken for the k-means++ algorithm to converge = ', end_time - start_time)
111 return clusters
112

```

Fig 12: Code for K-means++ algorithm

The convergence time of k-means is 6 seconds for k=3 without coordinates and that for k-means++ is 5 seconds. Thus k-means++ is faster than k-means.

Task-3: Face Recognition using Eigenface. (10 marks)

1. The script “my_images.py” converts image to gray scale, resizes them to standard size and stores them in folder “my_images”. The 10 images are labelled starting with “subject16” and of them one is test image called “subject16.test.jpg” being added to “test” folder and the rest to “train” folder of “my_images” directory.

Image alignment is important as it improves the accuracy of face recognition using Eigenface. The images are directly averaged and eigenvectors based on that average are used for eigenface. Hence, misalignment of image and background variations can lead to incorrect results.

2. Train an Eigen-face recognition system. Specifically, at least your face recognition system should be able to complete the following tasks:

(1) All the 135 training images from Yale-Face are read and each image is represented as a single data point in a high dimensional space. All the data points into a big data matrix.

(2) PCA has been performed on the data matrix X to find top k eigenvectors in the function pca() in “pca.py” (Fig 13). First the data (with each column as image) is normalized by subtracting mean from image. Then its covariance matrix is found. Next, eigenvectors and eigenvalues of the covariance matrix are calculated. The top k of Eigenvectors and Eigenvalues are returned with mean.

```
35
36 '''
37 pca function: Performs pca on matrix X and returns top k Eigenvalues, Eigenvectors, and mean of matrix X
38 '''
39 def pca(X, k):
40     mean = X.mean(axis=0)
41     # Subtract mean from images or datamatrix X
42     X = X - mean
43
44     n, m = X.shape
45     if n > m: # Use transpose trick when one of the dimensions of matrix is greater than other
46         C = X.T @ X # Covariance Matrix
47         eigval, eigvec = np.linalg.eigh(C)
48         eigvec = X @ eigvec
49         for i in range(m):
50             eigvec[:, i] = eigvec[:, i] / np.linalg.norm(eigvec[:, i])
51     else:
52         C = X @ X.T # Covariance Matrix
53         eigval, eigvec = np.linalg.eigh(C)
54
55     # Sort eigenvectors descending by their eigenvalue
56     idx = np.argsort(-eigval)
57     eigval = eigval[idx]
58     eigvec = eigvec[:, idx]
59     # select only top k
60     eigval = eigval[0: k].copy()
61     eigvec = eigvec[:, 0: k].copy()
62
63     return eigval, eigvec, mean
64
```

Fig 13: The function pca() performs PCA dimensionality reduction

The mean face obtained by averaging all the faces in training set is shown in image below. (Fig 14)



Fig 14: Mean face of all training images

Let a matrix S of dimension $n \times k$ be such that $k < n$. Then the non-zero eigenvalues of matrix SS^T are equal to the eigenvalues of matrix S^TS . And if the eigenvector of S^TS is v with eigenvalue λ , then the eigenvector of SS^T is Sv with same eigenvalue λ . It is then normalized (Theorem 1.4 from Week 06 Lect 09 Face at page 59).

The reason why this method works faster is that the matrix S^TS is a smaller $k \times k$ matrix compared to the larger SS^T matrix of $n \times n$ dimension. Using this theorem, we can find the eigenvectors of the large data matrix more quickly. This has been used in the code (Fig 13 – code lines: 44 to 50).

(3) The top 10 eigenfaces are plotted in image below. (Fig 15)

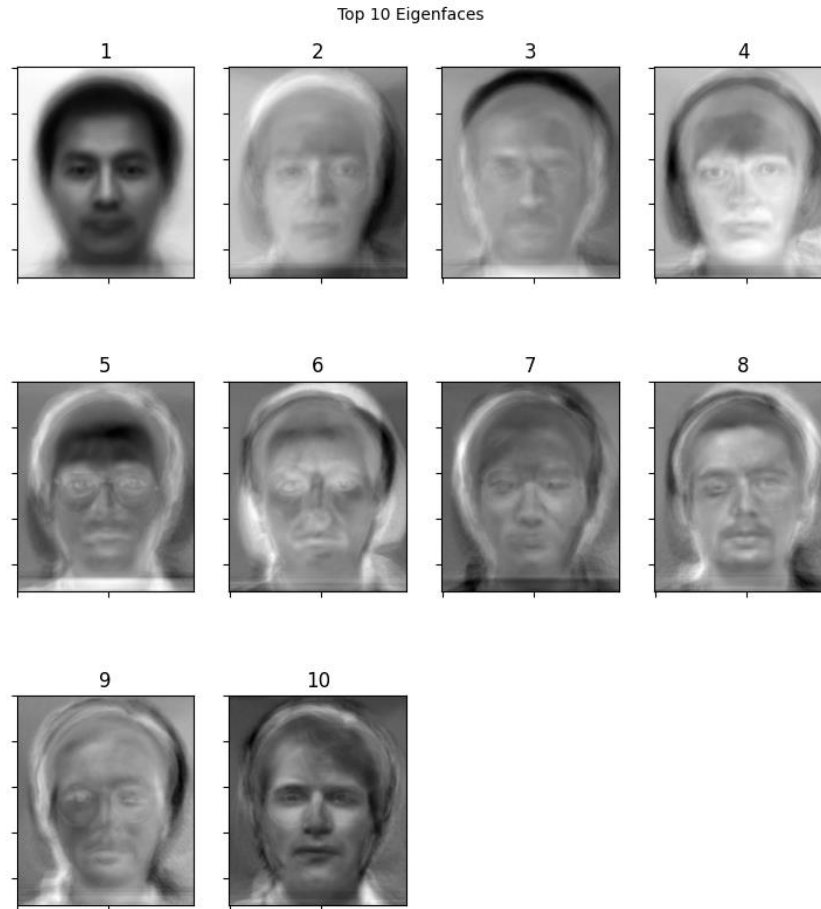


Fig 15: Top k ($k=10$) Eigenfaces

(4) For each of the 10 test images in the Yale-Face dataset, their projection onto the basis spanned by the top k eigenfaces is determined. This projection is used as a feature to perform nearest-neighbour search over all 135 faces and find out the top three face images that are most similar to the reference one are found. These top 3 faces are shown below in Table 2.

Table 2: Top 3 faces detected (2,3,4) to be closest to test face (1)

Subject 01

First 3 faces predicted for subject: 1



Subject 02

First 3 faces predicted for subject: 2



Subject 03

First 3 faces predicted for subject: 3



Subject 04

First 3 faces predicted for subject: 4



Subject 05

First 3 faces predicted for subject: 5



Subject 06

First 3 faces predicted for subject: 6



Subject 07

First 3 faces predicted for subject: 7



Subject 08

First 3 faces predicted for subject: 8



Subject 09

First 3 faces predicted for subject: 9



Subject 10

First 3 faces predicted for subject: 10



Accuracy is described in table below. (Table 3)

Table 3: Top-1 and Top-3 accuracies are listed below.

Accuracy (top-1)	Accuracy (top-3)
10%	30%

This is very poor accuracy. One possible reason for such low accuracy is the difference in lighting conditions in image and these images have varying background lighting.

(1) The result of running my test image through my face recognition system is in Fig 16.

First 3 faces predicted for my test image




Fig 16: Top 3 faces resulting from comparison of my test face with training set images

(2) The previous experiment is repeated by pre-adding the other 9 additional images of my face into the training set (a total of 144 training images). The top 3 faces that are the closest to my test face are displayed below (Fig 17).

First 3 faces predicted for my test image after training on my images as well



Fig 17: Top 3 faces resulting from comparison of my test face with the training set images including my image